Peer Review 1:

Analyzing Negation, Antonyms, and Tenses
For More Robust Natural Language Inference Models

Scope:

- **Rating:** Excellent
- **Justification:** Very thorough work that goes beyond basic model evaluation. Displays depth through creating custom contrast sets and adversarial examples targeting specific linguistic phenomena (negation, antonyms, tenses) and then fine-tuning the model to improve robustness.

Implementation:

- **Rating:** Excellent
- **Justification:** The author not only fine-tuned the ELECTRA model, but also made custom contrast sets and meticulously constructed adversarial training examples to address specific linguistic weaknesses. Introducing these new examples back into the training process, they improved the model's robustness. The steps are clearly detailed and appear technically sound, indicating a strong, ambitious implementation that successfully enhances the model's performance in a meaningful way, as well as demonstrating the utility of their methods.

Results/Analysis:

- **Rating:** Excellent
- **Justification:** The results and analysis clearly show the initial gaps in the model and dataset and the subsequent improvements achieved through their methodology. The author presented strong evidence and detailed insight, giving readers a compelling, fine-grained understanding of the work.

Clarity/Writing:

- **Rating:** Excellent
- **Justification:** The paper is clearly organized, following a nice logical flow. The writing is concise and coherent.